

Surface Electromyography in Silent Speech Interfaces: A Comprehensive Analysis of System Architectures, Form Factors, and Efficacy

Executive Summary

The evolution of Silent Speech Interfaces (SSIs) represents a profound paradigm shift in human-machine interaction, moving communication away from an exclusive reliance on acoustic wave generation toward the direct interpretation of neuromuscular activity. By leveraging surface electromyography (sEMG), these systems capture the efferent signal signatures transmitted from the central nervous system to the peripheral speech articulators, successfully intercepting motor intent even in the complete absence of vocalized sound. This report delivers an exhaustive, in-depth comparative analysis of six landmark studies that chart the developmental trajectory of sEMG-based speech recognition. These foundational benchmarks include Jorgensen's pioneering 2004 NASA Ames research on subauditory speech, Kapur's 2018 AlterEgo internal duplex architecture from the MIT Media Lab, Meltzner's 2018 clinical laryngectomy continuous speech application developed with Delsys and the NIH, the extensive 120-electrode high-density mapping investigations, Wu's 2024 overt-speech textile necklace form factor, and the ultra-low-power 2026 ETH Zurich SilentWear edge-computing platform.

Through a rigorous synthesis of these specific systems, several critical points of convergence and persistent architectural bottlenecks become readily apparent. The discipline demonstrates a powerful anatomical convergence toward the anterior neck and infrahyoid regions as the optimal recording sites, facilitating a crucial hardware transition from highly stigmatizing, clinical-grade facial electrode arrays to discreet, user-friendly textile neckbands. However, severe disagreements and profound technological gaps remain regarding the underlying electrode chemistry—specifically the viability of dry capacitive interfaces versus wet silver/silver chloride (Ag/AgCl) conductive gels. Furthermore, researchers continue to grapple with the stark performance degradation observed between overt, vocalized speech and purely silent, internal articulation, as well as the inherent trade-off between scaling vocabulary size and maintaining device wearability. This analysis uncovers deep, second- and third-order insights concerning the "wearability-performance paradox" and the pervasive, field-limiting challenge of inter-session electrode variability, ultimately outlining the strategic trajectory of neuroprosthetic communication toward multimodal sensor fusion and self-supervised foundation models.

Physiological Foundations of Electromyographic Speech Decoding

To properly contextualize the engineering challenges inherent in silent speech recognition, it is

essential to first outline the physiological and neurobiological foundations of speech production. Human vocal communication is an extraordinarily complex motor task that spans higher-order neural planning in the cerebral cortex, precise motor control signal transmission, localized muscle activation across multiple interconnected structures, and dynamic articulatory kinematics¹. When an individual intends to speak, the central nervous system transmits electrical action potentials down efferent nerve pathways to the neuromuscular junctions of over one hundred distinct muscles located across the face, jaw, neck, and larynx².

As acetylcholine bridges the synaptic cleft, it triggers the depolarization of the muscle fiber sarcolemma, generating motor unit action potentials (MUAPs). As these potentials propagate along the muscle fibers, their superposition creates a complex, spatiotemporal electrical field that permeates through the subcutaneous tissue, fat, and epidermis to the skin surface⁵.

Surface electromyography operates by detecting these microvolt-level potential differences across the skin. Critically for the field of silent speech interfaces, these neuromuscular signals are physically present regardless of the acoustic output. The electrical innervation occurs whether the speech is overtly vocalized with full glottal excitation, subvocally mouthed with normal articulatory movements but no airflow, or internally articulated in a purely covert, "mentally rehearsed" manner with microscopic or visually indiscernible muscle twitches⁸.

Because the sEMG signal represents the raw biological instruction set for speech and typically precedes actual physical articulatory motion by approximately 60 milliseconds, it offers a highly attractive, low-latency pathway for decoding intent¹². This mechanism is entirely immune to ambient acoustic noise, making it highly applicable for both specialized environments—such as aerospace operations, military covert communications, and loud industrial settings—and clinical interventions for individuals who have lost the ability to generate acoustic signals due to total laryngectomy, severe dysarthria, or neurological trauma¹⁴. However, capturing these signals presents a monumental engineering challenge. The sEMG signal is inherently non-stationary and highly susceptible to contamination from baseline drift, power-line interference, cardiac electrical artifacts, and shifting electrode-skin impedance⁴. Consequently, the history of sEMG-based speech decoding is defined by a relentless drive to isolate these fragile speech-specific potentials from a chaotic biological noise floor, balancing the need for pristine signal fidelity against the requirement for unobtrusive, wearable hardware.

Comprehensive Comparative Analysis of Benchmark Systems

To understand the current state, limitations, and future trajectory of silent speech recognition, a meticulous deconstruction of the specific methodologies, hardware configurations, vocabulary scales, and evaluation metrics of the field's most influential studies is required.

NASA Ames Subvocal Speech Recognition (Jorgensen, 2004)

Originating as a strategic initiative to aid astronauts and first responders operating in exceedingly noisy, microgravity, or low-pressure environments, the NASA Ames research program spearheaded by Chuck Jorgensen served as the foundational proof-of-concept for subauditory speech classification¹⁴. Relying on classical clinical electrophysiology techniques,

the hardware architecture utilized standard silver/silver chloride (Ag/AgCl) wet stick-on sensors¹⁴. The team deployed a minimalist, highly targeted array consisting of merely two electrode pairs—yielding a total of four electrodes—strategically placed on the anterior throat near the larynx and under the chin in the sublingual area⁹.

The NASA system operated on a highly restricted, closed-command vocabulary of 15 to 25 words, predominantly focusing on a core six-word navigation set designed to control a simulated Mars rover (stop, go, left, right, alpha, omega) alongside the ten English digits⁹. The signal processing pipeline involved amplifying the raw potentials between 20 and 500 Hertz, rectifying the waveforms, and applying complex dual-tree wavelet transforms (DTWT) or windowed short-time Fourier transforms (STFT)¹⁴. These spectral coefficients were then fed into a neural network pattern classifier¹⁴.

Under these heavily constrained conditions, the system achieved a 92.1% accuracy rate for overtly vocalized speech¹⁰. A subsequent systematic comparison between speaking modes revealed a crucial physiological performance gap: while vocalized speech yielded high accuracy, subvocally mouthed speech experienced signal attenuation, resulting in an 86.7% classification rate¹⁰. Most revealingly, early attempts by the NASA team to classify purely "mentally rehearsed" speech—where the subject merely thought the word without innervating the peripheral musculature—failed entirely. The researchers documented an absolute inability to conduct recognition in this mode due to a consistent lack of discernible sEMG activity at the skin surface, rendering performance no better than random chance⁷.

Calibration time for the NASA system initially required hundreds of labeled examples, encompassing several hours of intensive training data collection¹⁶. Through algorithmic optimization, the researchers eventually reduced the learning requirement to fewer than fifty examples per word, allowing a subject to train the neural network in the morning and utilize the interface by the afternoon¹⁶. The primary limitations identified by the Ames team included the total inability to scale to open-vocabulary or continuous speech, susceptibility to minor variations in electrode placement across sessions, and an absolute reliance on individual speaker dependency, requiring distinct models for every user¹⁴.

Internal Cognitive Coalescence: The AlterEgo Architecture (Kapur, 2018)

Developed a decade and a half later at the MIT Media Lab, the AlterEgo system sought to transcend basic command-and-control interfaces by melding human and machine cognition through an "internal," non-invasive peripheral nerve interface⁸. The hardware materialized as a highly distinctive, 3D-printed plastic wearable headset that hooked around the user's ear and extended down along the jawline, purposefully designed to avoid obstructing the user's mouth or vocal tract, thereby allowing simultaneous normal conversation⁸.

Kapur's architecture utilized 4 to 8 signal electrodes—specifically employing wet, gold-plated contacts that necessitated the application of conductive paste to lower skin impedance—positioned along the side of the face, cheek, and lower jaw⁴. The AlterEgo system was expressly optimized for silent, internal verbalization. Kapur characterized this internal

speech as the subtle, voluntary innervation of internal speech articulators that occurs when reading silently or talking to oneself, completely avoiding deliberate lip syncing or any visually discernible facial muscle movement³. The system also featured a closed-loop auditory feedback mechanism via bone conduction headphones, allowing the artificial intelligence agent to respond to the user silently³.

When evaluated on a highly restricted 20-word vocabulary tailored for specialized analytical tasks, such as arithmetic computation and chess commands, the underlying deep Convolutional Neural Network (CNN) achieved a median word transcription accuracy of 92%⁸. Calibration was highly streamlined compared to earlier systems, requiring only approximately 15 minutes of user-specific customization per session, followed by 90 minutes of execution trials⁸. Despite its high accuracy and conceptual brilliance, AlterEgo faced significant practical limitations. The physical plastic apparatus was bulky, highly visible, and socially stigmatizing²⁸. Furthermore, its reliance on wet electrode interfaces meant that signal fidelity would inevitably degrade as the conductive paste dehydrated over extended periods of use. Algorithmically, its performance was strictly confined to a minuscule, closed-command vocabulary, and it remained highly dependent on individual user calibration prior to every use²⁶.

Clinical Subvocal Decoding for Laryngectomy (Meltzner et al., 2018)

While NASA focused on aerospace applications and MIT focused on cognitive augmentation, Meltzner and colleagues at Delsys and the National Institutes of Health (NIH) approached sEMG decoding from a strictly clinical, neuroprosthetic perspective. Their primary objective was to restore naturalistic communication for patients who had undergone a total surgical laryngectomy due to trauma or disease¹⁵. This research represents a monumental benchmark in vocabulary scale, moving the field away from discrete, closed-command classification and toward the holy grail of open-vocabulary continuous speech recognition.

The Meltzner system utilized a tethered array of 8 to 11 miniaturized, single-differential sEMG sensors, affixed with double-sided medical adhesive and conductive gel¹⁵. These sensors were placed systematically across the face, targeting the supralabial and infralabial regions, and across the neck, monitoring the submental and ventromedial regions¹⁵. Because laryngectomized subjects have lost their vocal folds and glottal excitation, the research inherently centered on subvocally mouthed speech, where the user goes through the physical motions of articulation without producing sound³².

To process this data, the researchers extracted specialized Mel-frequency cepstral coefficients (MFCCs) adapted to the unique spectral characteristics of electromyography⁷. Rejecting the whole-word classification approach used by Jorgensen and Kapur, Meltzner's team developed sophisticated phoneme-based Hidden Markov Models (HMMs) combined with natural language equivalence grammars⁷. This phonetic approach allowed the system to sequence sounds logically, successfully decoding a massive 2,500-word continuous speech vocabulary encompassing over 1,200 unique phrases¹⁵. The results were unprecedented: the system achieved a 91.1% word recognition rate (translating to an 8.9% Word Error Rate) in healthy speakers, and an exceptional 89.7% accuracy rate in laryngectomized speakers³⁰. An ablation

study within the research noted that reducing the sensor set from the full array down to a minimal set of 4 targeted locations resulted in an average Word Error Rate increase to 13.6%¹⁵. While the algorithmic achievement and vocabulary scaling were unparalleled, the hardware configuration posed severe usability limitations. The reliance on multiple wired, gel-based sensors adhered directly to the user's face and throat created a highly stigmatizing, clinical, and uncomfortable setup, rendering the system completely impractical for unassisted daily life outside of a laboratory or hospital environment².

High-Density sEMG Electrode Mapping (Zhu, Wang et al., 2020)

To systematically address the seemingly arbitrary nature of electrode placement in previous studies—where researchers relied on educated anatomical guesses to target specific muscles—researchers deployed high-density (HD) sEMG arrays to comprehensively map the entire neuro-electrical topography of the facial and neck articulatory network³⁴. Utilizing a clinical-grade 120-channel REFA system operating at a 2048 Hz sampling rate, the team applied matrices of wet, gel-injected electrodes covering almost the entirety of the cheeks, chin, and neck³⁴. The study involved subjects overtly vocalizing ten English digits in addition to maintaining a silent baseline mode³⁴.

While the full 120-channel array yielded classification accuracies approaching 99% using robust time-domain features (such as waveform length and mean absolute value) fed into a Linear Discriminant Analysis (LDA) classifier, the primary scientific contribution of the study was its rigorous dimensionality reduction¹⁸. By applying a Sequential Forward Selection (SFS) algorithm, the researchers proved mathematically that the vast majority of the 120 electrodes were entirely redundant. A highly optimized subset of just 4 optimally placed channels could recover an 80% classification accuracy, while a carefully selected configuration of 10 to 20 channels was sufficient to achieve 90% to 95% accuracy³⁴.

This study exposed the critical limitations of high-density arrays for practical speech interfaces: massive channel counts introduce overwhelming computational processing burdens, high signal cross-talk, tedious maintenance and preparation expenses, and extreme physical discomfort for the subjects due to dense wiring³⁵. However, the research established a definitive physiological baseline, proving that highly accurate speech recognition does not require dense, full-facial coverage, provided the remaining electrodes are strategically positioned over high-yield motor units.

The Textile Necklace Form Factor (Wu et al., 2024)

Addressing the social stigmatization, skin irritation, and general discomfort of facial electrodes highlighted by previous generations of hardware, Wu and colleagues introduced a paradigm-shifting form factor: a discreet, wireless, textile-based neckband¹⁹. This system sought to entirely eliminate facial electrode placements, relying instead on 10 reusable, 3D-printed, electroless gold-plated dry electrodes distributed evenly around an elastic collar matching the circumference of the human neck¹⁹. Signal acquisition was handled by an ultra-low-profile WANDmini recording module, eliminating tethered data acquisition systems¹⁹.

The researchers sought to decode an 11-word vocabulary consisting of highly similar articulatory phonological structures (e.g., "aga", "aba") across multiple speakers³⁶. The neckband achieved a striking 92.7% user-dependent classification accuracy using Random Forest and Deep Learning architectures, proving that the musculature of the neck alone—capturing infrahyoid, suprahyoid, and laryngeal activity—contains sufficient signal fidelity for discrete word classification¹⁹.

However, a profound limitation of the 2024 study lies in its evaluation methodology: the impressive 92.7% accuracy metric was achieved exclusively using overt, vocalized speech data³⁶. The authors explicitly noted that decoding overt speech heavily benefits from secondary physical phenomena—such as active glottal vibration, gross tissue displacement, and airflow dynamics—which are significantly attenuated or entirely absent during purely silent speech production³⁶. Furthermore, specific ablation studies within the research revealed that reducing the sensor count from 10 down to only two electrodes—a configuration commonly used in simple electroglottographs—resulted in a catastrophic degradation of classification accuracy, proving that circumferential spatial coverage around the neck remains absolutely necessary when utilizing noisy, dry electrodes¹⁹.

Ultra-Low-Power Edge Processing: SilentWear (Spacone, Frey et al., ETH Zurich, 2026)

Representing the apex of current hardware engineering in wearable SSIs, the ETH Zurich SilentWear system elegantly synthesized advanced textile design with embedded, ultra-low-power edge computing²⁸. Building upon the conceptual framework of the Wu architecture, SilentWear utilized a fully dry, textile-based neckband constructed with graphene and PEDOT:PSS-coated fibers, but expanded the sensor array to 14 fully differential EMG channels to maximize spatial sampling¹³. The hardware was powered by the specialized BioGAP-Ultra platform, which facilitated on-device, real-time neural network inference consuming merely 20.5 milliwatts of total power, yielding over 27 hours of continuous battery life²⁸.

The ETH system evaluated an 8-command Human-Machine Interaction (HMI) vocabulary (including navigational words like "up," "down," "start," "stop," plus a rest class) using a highly efficient 15,000-parameter CNN dubbed "SpeechNet"²⁸. Crucially, unlike the Wu study, the Swiss researchers conducted a direct, rigorous, within-subjects comparison between overtly vocalized and purely silent speech under identical hardware and algorithmic conditions. In a within-session, cross-validated environment, vocalized speech yielded an 84.8% accuracy, while the accuracy for purely silent speech dropped to 77.5%²⁸.

The most scientifically significant contribution of the SilentWear study, however, was its unflinching, quantified analysis of "inter-session variability." When the neckband was removed by the user and repositioned on a subsequent day, the "leave-one-session-out" accuracy plummeted to 71.1% for vocalized speech and a mere 59.3% for silent speech¹³. This demonstrated that micro-shifts in dry electrode positioning over the underlying muscle bellies completely shatter the latent feature maps learned by the neural network. To combat this severe

degradation, the researchers introduced an incremental fine-tuning strategy. They demonstrated that gathering less than 10 minutes of new calibration data upon donning the device could successfully update the neural network weights and recover more than 10% of the lost classification accuracy²⁸.

Structural Comparison Matrix

Syst em / Stud y	Voca bular y Size	Best Repo rted Accu racy / WER	Elect rode Coun t	Plac eme nt Strat egy	Form Fact or	Dry vs. Gel	Silen t vs. Voca lized Eval uatio n	Calib ratio n / Train ing Time	Key State d Limit ation s
Jorg ense n (2004) [cite: 9, 10, 14, 16]	15-25 word s (6 com mand s + 10 digits)	92.1 % (Voca lized) , 86.7 % (Mout hed)	4 (2 speci fic pairs)	Anter ior throat (laryn x), Subli ngual chin	Wire d stick- on sens ors	Wet (Ag/A gCl)	Expli cit comp ariso n: perfo rman ce drops for mout hed spee ch; fails on intern al thoug ht.	Initiall y sever al hours ; optim ized to <50 iterati ons per word.	Strict singl e-spe aker depe nden cy; total failur e on purel y ment al rehea rsal; close d voca bular y.
Kapu r (2018) [cite:	20 word s (Arith metic	92% (Medi an word accur	4 to 8	Face, Chee k, Jawli ne	Plasti c rigid face/ neck	Wet (Gold -plate d with	Focu ses entire ly on intern	15 minut es of user custo	Stigm atizin g aesth etic

8, 11, 26, 28]	& Ches s focus)	acy)			head set	cond uctiv e paste)	al/sile nt spee ch, omitti ng direct vocali zed comp ariso n.	mizat ion per sessi on.	form factor ; minu scule close d-co mma nd voca bular y; requir es wet paste .
Meltz ner (2018) [cite: 15, 17, 30]	2,500 word s (Cont inuou s sequ ence s)	91.1 % Healt hy (8.9 % WER), 89.7 % Clinic al	8 to 11	Supr alabi al, Infral abial, Sub ment al, Ventr al Neck	Minia turize d tether ed array	Wet (Gel/ Skin prep requir ed)	Focu ses heavil y on subv ocal ly mout hed spee ch for laryn gecto my patie nts.	Inten sive HMM phon eme sequ ence traini ng requir ed.	Highl y obtru sive facial place ment; unus able in indepe nde nt daily civilia n life.
Zhu/ Wan g (2020) [cite:	11 tasks (10 Engli sh digits	>95% (120 ch), 80% (opti mize	120	Full topog raphi cal mapp ing of	HD sEM G Matri x Array	Wet (Con ducti ve gel inject	Analy zed overt vocali zed spee	28 repeti tions per digit recor	Extre me comp utatio nal proce

18, 34, 35]	+ silent rest mode)	d 4 ch)		Face and Neck		ed)	ch against a silent baseline state.	ded.	ssing load; severe physical wire discomfort; restricted to laboratory settings.
Wu (2024) [cite: 19, 36, 38]	11 words (highly similar phonology, e.g., "aga")	92.7 %	10	Circumferential Neck	Textile wireless Necklace	Dry (Gold-plated 3D printed)	Explicitly relies on overt/vocalized speech; notes silent speech lacks mechanical cues.	Multi-speaker data pooling utilized.	Accuracy claim relies entirely on overt vocalized signals; highly limited vocabulary scale.
SilentWearer	8 commands	84.8 % (Voca	14	Circumferential	Textile wireless	Dry (Graphene/	Explicit comp	<10 minutes	Severe inter-

(2026) [cite: 28, 39, 40, 42]	s + 1 rest class	lized), 77.5% (Silent)		Neck	ess Neck band	PED OT:PSS)	arison: silent speech incur s a ~7% accuracy penalty vs vocalized.	required for incremental fine-tuning recovery.	session accuracy degradation due to repositioning shift; restricted vocabulary.
----------------------------------	------------------	------------------------	--	------	---------------	-------------	--	--	---

Points of Convergence and Architectural Consensus

Across over two decades of concentrated research spanning aerospace engineering, cognitive science, and clinical neuroprosthetics, the scientific literature demonstrates several powerful points of convergence. These shared methodologies signal the maturation of sEMG-based speech decoding from a theoretical laboratory novelty to a highly applied neuroprosthetic discipline.

The Anterior Neck and Infrahyoid Signal Viability

Early assumptions in the field of speech decoding hypothesized that highly accurate classification fundamentally required direct electrical surveillance of the highly dynamic facial muscles—specifically targeting the levator anguli oris, the zygomaticus major, the orbicularis oris, and the broader perioral musculature responsible for labial shaping³¹. Consequently, both Kapur’s AlterEgo and Meltzner’s clinical array relied heavily on obtrusive facial placements to capture the exact nuances of lip, cheek, and jaw movements⁸.

However, subsequent empirical evidence has converged overwhelmingly on the anterior and lateral neck as an equally viable, and infinitely more practical, signal source. The neck houses a dense network of muscles integral to speech, including the sternocleidomastoid, the platysma, and the complex web of infrahyoid and suprahyoid muscles (such as the digastric and mylohyoid) that stabilize the larynx and forcefully manipulate the base of the tongue during phonation⁴³. The clinical success of the Wu 2024 necklace (achieving 92.7% accuracy) and the SilentWear 2026 platform (achieving up to 84.8% accuracy) relying strictly on neck-based dry electrodes proves definitively that the neck region provides a sufficiently high Signal-to-Noise

Ratio (SNR) for robust algorithmic classification without requiring facial telemetry¹⁹.

The Shift Toward the Neckband Form Factor

Directly downstream of the anatomical consensus regarding the infrahyoid musculature is a profound convergence on hardware design: the transition away from tethered facial arrays toward the completely wireless, circumferential neckband¹⁹. High-density mapping studies unequivocally proved that attaching arrays of 120 electrodes across the cheeks and jaw is physiologically intolerable for sustained periods, inducing muscular fatigue, skin irritation from adhesives, and severe social stigma²⁸. By migrating the entire sensor payload to a fabric-based textile neckband, modern researchers have achieved a form factor that closely mimics conventional, socially accepted consumer electronics, such as audio headphones or wearable fitness monitors³⁹. This consensus marks the definitive transition of SSIs from clinical diagnostic tools to viable ambient computing interfaces.

Algorithmic Evolution and Edge Processing Dominance

There is an unmistakable, linear progression in the mathematical complexity and deployment strategy of the underlying machine learning architectures used to decode speech intent. Early systems, such as Jorgensen's and the HD mapping studies, relied heavily on labor-intensive, manual feature extraction—calculating waveform lengths, zero-crossing rates, and localized mean absolute values—which were then fed into shallow classifiers like Linear Discriminant Analysis (LDA) or standard Support Vector Machines (SVM)³⁵. Meltzner significantly advanced this paradigm by utilizing biologically-inspired Mel-frequency cepstral coefficients (MFCCs) paired with Hidden Markov Models (HMMs) to map continuous sequences of phonemes over time⁷.

Modern systems, encompassing both the AlterEgo and SilentWear architectures, have almost universally converged on Deep Learning—specifically Deep Convolutional Neural Networks (CNNs) and sequence prediction algorithms utilizing Connectionist Temporal Classification (CTC)⁴. These deep networks autonomously extract hierarchical spatio-temporal features directly from the raw or mildly filtered sEMG time-series data, eliminating the need for handcrafted features. Furthermore, the SilentWear 2026 study highlights the arrival of edge-AI capability; successfully executing a 15,000-parameter CNN directly on a microcontroller (consuming only 63.9 μ J per inference) proves that these sophisticated algorithms no longer require tethered, external computational mainframes to function²⁸.

Divergence, Disagreements, and Technological Gaps

Despite the strong anatomical and algorithmic convergence, the discipline is fractured by several fundamental disagreements and glaring technological gaps that currently prevent the mass commercialization of silent speech devices.

The Dry Versus Gel Electrode Dichotomy

The most significant physical barrier in sEMG recording is the complex biophysics of the skin-electrode interface. Classical electrophysiology strictly dictates the use of silver/silver

chloride (Ag/AgCl) electrodes paired with a hydrophilic conductive gel or paste. This gel acts as a transducer, bridging the highly resistive outer layer of dead skin cells (the stratum corneum), effectively lowering contact impedance to below 5 k Ω to ensure high-fidelity signal capture¹⁸. Consequently, Jorgensen, Kapur, Meltzner, and the 120-electrode mapping study all utilized wet or paste-assisted electrodes to guarantee signal integrity¹⁸. However, wet electrodes are fundamentally flawed for continuous daily use: the conductive gel dehydrates over time leading to signal loss, causes dermatological irritation, and is entirely unsuited for a "don-and-forget" consumer wearable¹³.

Conversely, modern textile architectures like Wu (2024) and SilentWear (2026) forcefully advocate for completely dry, reusable electrodes, utilizing gold-plated polymers or functionalized graphene/PEDOT:PSS textiles¹⁹. While dry electrodes radically improve user comfort and device longevity, they introduce severe capacitive interference, extreme susceptibility to motion artifacts, and dramatically higher baseline impedances (e.g., a mean impedance of 55.1 k Ω reported by Wu)¹⁹. The literature remains deeply split: clinical researchers demanding high-accuracy, continuous speech models continue to rely on wet adhesion to maintain signal fidelity, whereas HCI and consumer-focused engineering labs have accepted the degraded SNR and higher noise floor of dry electrodes as a necessary, unavoidable compromise for everyday wearability.

The Persistent Gap Between Vocalized and Silent Speech

A critical semantic and physiological divide exists regarding the exact nature of the speech being decoded by these interfaces. When a user speaks aloud (overtly/vocalized), the sEMG signal is massively bolstered by vigorous muscle contractions, significant physical articulatory displacement, resonant glottal vibrations, and localized changes in skin tension driven by expelled airflow³⁶. Conversely, when a user shifts to silently mouthed speech, or entirely internal covert subvocalization, the physical excursion of the muscles decreases exponentially, resulting in a microscopic millivolt signal amplitude¹¹.

This unalterable physiological reality manifests as a consistent, severe accuracy penalty across the literature. As early as 2004, Jorgensen noted a distinct drop from 92.1% accuracy for vocalized speech down to 86.7% for mouthed speech¹⁰. Most tellingly, the highly optimized 2026 SilentWear study documented an 84.8% accuracy for vocalized commands, which fell precipitously to 77.5% when the exact same vocabulary was performed silently²⁸. Some recent studies, such as the Wu 2024 necklace, claim exceptionally high accuracies (92.7%) but explicitly restrict their primary evaluation to overt, vocalized speech classifiers, openly acknowledging that decoding purely silent speech is drastically more difficult due to the total attenuation of mechanical acoustic cues³⁶. Bridging this classification gap between full physical articulation and nearly imperceptible subvocalization remains a major, unsolved hurdle in the pursuit of true silent communication.

Electrode Density Versus Practical Usability

While the 120-electrode high-density mapping study proved mathematically that optimal subset selection could reduce required channels down to a highly efficient 4 or 10, the broader question

of minimum spatial coverage remains highly contested³⁴. Jorgensen utilized only 4 electrodes to achieve his results, whereas Kapur used up to 8⁴. However, as newer systems migrated to the neck and adopted inherently noisier dry interfaces, sensor counts began climbing again to compensate for the loss of rich facial data and the increased electrical noise floor. Wu's necklace utilized 10 electrodes, and SilentWear required 14 fully differential channels¹⁹. There is currently no consensus on the minimum necessary sensor topology required to decode a broad vocabulary using dry electrodes. The ablation study in Wu's 2024 paper specifically noted that restricting the system to only two electrodes—a configuration commonly seen in basic electroglottographs—resulted in catastrophic accuracy degradation, proving that widespread, dense spatial sampling around the neck circumference is still absolutely required to extract intelligible features¹⁹.

Vocabulary Constraints and the Scalability Wall

Perhaps the most glaring limitation across the modern, wearable SSI landscape is the dramatic collapse of vocabulary size. In 2018, using tethered facial electrodes and phoneme-based sequence modeling, Meltzner successfully decoded a massive 2,500-word continuous vocabulary¹⁵. Yet, as the field advanced toward highly desirable, consumer-friendly dry neckbands (AlterEgo, Wu 2024, SilentWear 2026), vocabularies plummeted to 20, 11, and 8 discrete words, respectively¹⁹.

This severe regression exposes a fundamental limitation in current machine learning classification paradigms. Treating SSI as a closed-set classification problem—where a neural network learns a static, distinct mapping for a specific whole word—cannot scale to the tens of thousands of words required for fluid, natural human conversation⁴³. Without sufficient spatial resolution and signal-to-noise ratio to reliably isolate individual sub-word phonemes and track complex linguistic coarticulation dynamics, dry neckbands are currently trapped in the limited domain of simple, isolated control commands (e.g., "up," "down," "start," "stop")²⁸. Breaking through this scalability wall represents the ultimate threshold for the technology.

Second and Third-Order Insights

Beyond the direct empirical results of accuracy percentages and channel counts, analyzing the chronologic and architectural shifts between these six benchmark studies reveals deeper, underlying paradigms governing the intersection of human physiology and artificial intelligence.

The Wearability-Performance Paradox

The trajectory of sEMG speech recognition exhibits a distinct, punishing paradox: as the physical hardware interface becomes more human-centric and socially acceptable, the biological signal becomes increasingly hostile to machine interpretation. By stripping away wet conductive gels and migrating sensors away from the highly expressive, easily isolatable facial muscles to a discreet, consolidated neck collar, engineers drastically reduce social stigmatization²⁸. However, the direct cost of this aesthetic improvement is the introduction of massive signal artifacts, extreme capacitive electrical noise, cross-talk from adjacent non-speech musculature (such as the muscles used for swallowing or basic head rotation), and

severe amplitude attenuation of the electrical potentials³⁷.

Consequently, to maintain the illusion of high system accuracy (i.e., keeping performance metrics artificially above 90% for publication), researchers are often forced to constrain the operating environment by shrinking the vocabulary back to a handful of highly distinct discrete commands. True cognitive coalescence—as ambitiously envisioned by the AlterEgo architecture—cannot occur until the algorithmic backend becomes robust enough to extract delicate, phoneme-level continuous speech data from the highly degraded, noisy signal provided by a comfortable dry neckband⁸.

The Tyranny of Inter-Session Variability

The most profound, existential barrier to the real-world deployment of SSIs is not achieving initial accuracy, but rather maintaining temporal stability, frequently referred to in the literature as inter-session variability²⁸. In highly controlled laboratory settings, an electrode array is carefully attached, the system is calibrated to the user, and the testing is conducted immediately. Under these pristine conditions, the spatial relationship between the metal electrode and the underlying motor unit action potentials remains perfectly static.

However, in real-world scenarios, a consumer will remove the neckband at the end of the day and put it back on the following morning. A microscopic shift of even a few millimeters completely changes the relative amplitude, phase, and morphology of the sEMG signal captured by the differential channels¹³. The 2026 SilentWear study explicitly quantified this catastrophic failure mode: a model meticulously trained to 84.8% accuracy plummeted to 71.1% simply because the user took the device off and put it back on²⁸. To combat this, the researchers were forced to introduce an incremental fine-tuning strategy, requiring the user to spend up to 10 minutes recalibrating the device before every single use to recover baseline performance²⁸. This reveals a critical third-order insight: static machine learning models, no matter how deep or complex, are fundamentally incompatible with the chaotic reality of wearable electrophysiology. Future systems must employ continuous, unsupervised domain adaptation or self-supervised learning algorithms that constantly update their internal synaptic weights in the background to account for micro-shifts in electrode positioning.

Neuromuscular Redundancy and Phonological Confusion

The SFS dimension-reduction demonstrated in the 120-electrode mapping study³⁴, evaluated alongside the deep phonological confusion analyses presented in the Wu 2024 paper¹⁹, highlights the vital concept of neuromuscular redundancy within the human body. Because the facial and neck muscles are intricately interconnected—largely lacking rigid bony insertions and acting instead as a complex, overlapping biological web—the cognitive intent to articulate a specific phoneme sends a ripple of stabilizing contractions and electrical activations across the entire craniofacial region⁵⁵.

Wu's research definitively proved that the machine learning classification confusions observed using only data from the neck perfectly mirror the linguistic confusions seen when using data gathered from the entire face¹⁹. This implies that the infrahyoid and laryngeal neck muscles act as a low-resolution, holographic "mirror" for the entire articulatory apparatus. The implication of

this finding is staggering: hardware engineers do not actually need to place highly obtrusive sensors directly over the lips or jaw to know exactly what the lips and jaw are doing. The neuromuscular stabilization patterns radiating through the neck contain a highly compressed, overlapping mathematical representation of the entire vocal tract's intent. The ultimate challenge lies entirely in designing neural network decoders sophisticated enough to untangle and uncompress this highly dense, convoluted signal.

Strategic Outlook and Future Trajectories

To shatter the persistent scalability wall and overcome the punishing wearability-performance paradox, the next generation of Silent Speech Interfaces must evolve beyond the rigid constraints of isolated sEMG classification.

1. **Multimodal Sensor Fusion:** Relying exclusively on capturing electrical potentials via high-impedance dry skin contact has reached a point of diminishing returns. The future of robust speech decoding lies in multi-modal hardware architectures. Integrating sEMG with alternative, non-electrical sensing modalities that are entirely immune to skin impedance changes—such as miniaturized inertial measurement units (IMUs) designed to capture the microscopic tissue vibrations of mechanical mechanomyography, or highly sensitive wearable magnetometers deployed to detect magnetomyographic (MMG) fields—can provide the orthogonal data streams necessary to disambiguate overlapping phonemes³⁶. By fusing electrical action potentials with direct mechanical and magnetic displacement data, edge-AI systems could finally differentiate between highly similar articulatory gestures without ever requiring the return of facial electrodes.
2. **Self-Supervised Foundation Models for Electrophysiology:** The current reliance on fully supervised machine learning—which painfully requires thousands of labeled speech repetitions per individual user—must be fundamentally abandoned. Inspired by the massive success of large language models (LLMs) in text generation, the neuroprosthetic field must pivot toward self-supervised foundation models pre-trained on massive, unlabeled, cross-population datasets of human neuromuscular activity¹². By forcing a neural network to learn the intrinsic, latent physical structure of human motor control prior to specific deployment, such generalized models could achieve zero-shot or few-shot adaptation to entirely new users. This architectural leap would drastically reduce the dreaded 15-minute calibration times and inherently mitigate the destructive effects of inter-session electrode shift, paving the way for true plug-and-play interfaces⁵⁰.

The pursuit of a seamless, non-invasive Silent Speech Interface represents one of the most ambitious and transformative undertakings in modern human-machine interaction. Analyzing the two-decade trajectory from Jorgensen's early throat-sensor experiments at NASA in 2004, through the pioneering conceptual frameworks of Kapur's AlterEgo and Meltzner's clinical phonetic decoding in 2018, to the highly refined, edge-processed textile neckbands of Wu and the ETH Zurich SilentWear team, reveals an extraordinary, relentless maturation of both hardware engineering and algorithmic design.

The industry has achieved an irrefutable, permanent consensus: the anterior neck provides a sufficient neuromuscular signal to bypass the social stigmatization of facial arrays, and the transition from wet conductive gels to dry textile electrodes is absolutely mandatory for

widespread consumer adoption. Yet, this physical liberation has exacted a heavy toll on raw signal fidelity, temporarily trapping wearable systems in closed-vocabulary, single-command paradigms, and exposing severe vulnerabilities to inter-session variability and electrode shifting. The path forward demands a radical departure from static classification models. By fully embracing continuous domain adaptation, multimodal sensor fusion, and self-supervised neuromuscular foundation models, researchers can bridge the fidelity gap, eventually realizing the profound vision of continuous, purely silent cognitive coalescence between human biological intent and digital computation.

Works cited

1. Towards EMG-to-Speech with Necklace Form Factor | Request PDF - ResearchGate, https://www.researchgate.net/publication/383652463_Towards_EMG-to-Speech_with_Necklace_Form_Factor
2. (PDF) Silent Speech Interfaces for Speech Restoration: A Review - ResearchGate, https://www.researchgate.net/publication/344153961_Silent_Speech_Interfaces_for_Speech_Restoration_A_Review
3. Arnav Kapur's AlterEgo Innovation | PDF | Artificial Intelligence - Scribd, <https://www.scribd.com/document/431644853/Arnav-Kapur-Fact-Sheet-PDF-Lemelson-Mit>
4. A Continuous Silent Speech Recognition System for AlterEgo, a Silent Speech Interface - Eric Wadkins, http://www.ericwadkins.com/doc/MEng_Thesis_FINAL.pdf
5. A thin, flexible multielectrode grid for high-density surface EMG, <https://journals.physiology.org/doi/10.1152/japplphysiol.00521.2003>
6. Neural decoding from surface high-density EMG signals: influence of anatomy and synchronization on the number of identified motor units | bioRxiv, <https://www.biorxiv.org/content/10.1101/2022.02.05.479100.full>
7. Development of sEMG sensors and algorithms for silent speech recognition - PMC - NIH, <https://pmc.ncbi.nlm.nih.gov/articles/PMC6168082/>
8. Humans and Computers Talk Silently With New MIT Headset - Labroots, <https://www.labroots.com/trending/technology/8464/humans-computers-talk-silently-mit-headset>
9. NASA Develops System To Computerize Silent, 'Subvocal Speech' - ScienceDaily, <https://www.sciencedaily.com/releases/2004/03/040318072412.htm>
10. Speech Recognition for Vocalized and Subvocal Modes of Production Using Surface EMG Signals from the Neck and Face - ResearchGate, https://www.researchgate.net/profile/Marcelo-Silva-12/post/I-want-to-acquire-speech-emg-signal-using-surface-electrodes-where-and-in-what-configuration-electrodes-should-be-placed/attachment/59d629d479197b8077988563/AS%3A337357501943811%401457443704296/download/c_010.pdf
11. AlterEgo: A Personalized Wearable Silent Speech Interface - Andy Matuschak, <https://andymatuschak.org/files/papers/Kapur%20et%20al%20-%202018%20-%2>

[0AlterEgo.pdf](#)

12. Advances in EMG-to-Speech Generation - Emergent Mind, <https://www.emergentmind.com/topics/emg-to-speech-generation>
13. A Parallel Ultra-Low Power Silent Speech Interface based on a Wearable, Fully-dry EMG Neckband - arXiv, <https://arxiv.org/pdf/2509.21964>
14. In the News - IEEE Computer Society, <https://www.computer.org/csdl/magazine/ex/2004/03/x3004/13rUwI5Uc8>
15. Silent Speech Recognition as an Alternative Communication Device for Persons with Laryngectomy - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC5851476/>
16. System susses out silent speech TRN 032404 - Technology Research News, http://www.trnmag.com/Stories/2004/032404/System_susses_out_silent_speech_032404.html
17. Silent Speech Recognition as an Alternative Communication Device for Persons with Laryngectomy - PubMed, <https://pubmed.ncbi.nlm.nih.gov/29552581/>
18. A Pilot Study on the Performance of Time-Domain Features in Speech Recognition based on high-density sEMG, <https://ieeexplore.ieee.org/iel7/9629355/9629471/09630541.pdf>
19. Towards EMG-to-Speech with a Necklace Form Factor - ISCA Archive, https://www.isca-archive.org/interspeech_2024/wu24k_interspeech.pdf
20. Throat Talkin' | Discover Magazine, <https://www.discovermagazine.com/throat-talkin-15380>
21. Improving Situational Awareness for First Responders via Mobile Computing, <https://ntrs.nasa.gov/api/citations/20060000029/downloads/20060000029.pdf>
22. (PDF) Web Browser Control Using EMG Based Sub Vocal Speech Recognition., https://www.researchgate.net/publication/221177953_Web_Browser_Control_Using_EMG_Based_Sub_Vocal_Speech_Recognition
23. Small-vocabulary speech recognition using surface electromyography | Request PDF, https://www.researchgate.net/publication/222822259_Small-vocabulary_speech_recognition_using_surface_electromyography
24. Human-machine cognitive coalescence through an internal duplex interface - DSpace@MIT, <http://dspace.mit.edu/handle/1721.1/120883>
25. MIT's new voiceless interface can read the words in your head | Next Nature, <https://nextnature.org/en/magazine/story/2018/alter-ego>
26. MIT Researchers Have Created a Bizarre Headset That Lets You Communicate Without Speaking - Science Alert, <https://www.sciencealert.com/silent-voice-headset-subvocalisation-computer-interface-mit>
27. MIT is making a device that can 'hear' the words you say silently - Popular Science, <https://www.popsci.com/device-hears-your-silent-speech/>
28. SilentWear: an Ultra-Low Power Wearable System for EMG-based Silent Speech Recognition - arXiv, <https://arxiv.org/html/2603.02847v1>
29. ebowwa/Open-Altergo - GitHub, <https://github.com/ebowwa/Open-Altergo>
30. Development of sEMG sensors and algorithms for silent speech recognition - PubMed - NIH, <https://pubmed.ncbi.nlm.nih.gov/29855428/>

31. Electrode Setup for Electromyography-Based Silent Speech Interfaces: A Pilot Study - MDPI, <https://www.mdpi.com/1424-8220/25/3/781>
32. Surface Electromyography–Based Recognition, Synthesis, and Perception of Prosodic Subvocal Speech - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC8740708/>
33. A Wearable Silent Text Input System Using EMG and Piezoelectric Sensors - MDPI, <https://www.mdpi.com/1424-8220/25/8/2624>
34. The Effects of Channel Number on Classification Performance for sEMG-based Speech Recognition - PubMed, <https://pubmed.ncbi.nlm.nih.gov/33018661/>
35. The Effects of Channel Number on Classification Performance for sEMG-based Speech Recognition - IEEE Xplore, <https://ieeexplore.ieee.org/iel7/9167168/9175149/09176260.pdf>
36. Silent Speech Recognition with Wearable Magnetometers - bioRxiv, <https://www.biorxiv.org/content/10.1101/2025.08.04.668236v1.full-text>
37. Silent Speech Recognition with Wearable Magnetometers - bioRxiv, <https://www.biorxiv.org/content/10.1101/2025.08.04.668236v2.full.pdf>
38. [2407.21345] Towards EMG-to-Speech with a Necklace Form Factor - arXiv, <https://arxiv.org/abs/2407.21345>
39. (PDF) SilentWear: an Ultra-Low Power Wearable System for EMG-based Silent Speech Recognition - ResearchGate, https://www.researchgate.net/publication/401522892_SilentWear_an_Ultra-Low_Power_Wearable_System_for_EMG-based_Silent_Speech_Recognition
40. [2603.02847] SilentWear: an Ultra-Low Power Wearable System for EMG-based Silent Speech Recognition - arXiv, <https://arxiv.org/abs/2603.02847>
41. (PDF) Wireless Silent Speech Interface Using Multi-Channel Textile EMG Sensors Integrated into Headphones - ResearchGate, https://www.researchgate.net/publication/393075385_Wireless_Silent_Speech_Interface_Using_Multi-Channel_Textile_EMG_Sensors_Integrated_into_Headphones
42. SilentWear: an Ultra-Low Power Wearable System for EMG-based Silent Speech Recognition - arXiv, <https://arxiv.org/html/2603.02847v2>
43. Silent Speech Recognition using Electromyography Signals - CSE - IIT Kanpur, https://www.cse.iitk.ac.in/users/amitangshu/Silent_Speech_Recognition_using_Electromyography_Signals.pdf
44. Silent Speech Decoding Using Spectrogram Features Based on Neuromuscular Activities, <https://www.mdpi.com/2076-3425/10/7/442>
45. Electrode Setup for Electromyography-Based Silent Speech Interfaces: A Pilot Study, https://www.researchgate.net/publication/388475742_Electrode_Setup_for_Electromyography-Based_Silent_Speech_Interfaces_A_Pilot_Study
46. Silent Speech Decoding Using Spectrogram Features Based on Neuromuscular Activities, <https://pmc.ncbi.nlm.nih.gov/articles/PMC7407985/>
47. ORIGINAL RESEARCH PAPER Anuj Kabra - Worldwidejournals.com, https://www.worldwidejournals.com/paripex/recent_issues_pdf/2025/March/surface-electromyography-for-silent-speech-recognition-applications-in-physiotherapy-a

- [nd-assistive-technologies_March_2025_2812629747_3206789.pdf](#)
48. 42nd Annual International Conference of the IEEE Engineering in Medicine & Biology Society, EMBC 2020, Montreal, QC, Canada, July 20-24, 2020 - ResearchGate, <https://researchgate.org/publication/embc-2020>
 49. AlterEgo: A Personalized Wearable Silent Speech Interface | Request PDF - ResearchGate, https://www.researchgate.net/publication/323669071_AlterEgo_A_Personalized_Wearable_Silent_Speech_Interface
 50. emg2speech: synthesizing speech from electromyography using self-supervised speech models - arXiv, <https://arxiv.org/html/2510.23969v1>
 51. PulpBio/SilentWear · Datasets at Hugging Face, <https://huggingface.co/datasets/PulpBio/SilentWear>
 52. Sebastian FREY | PhD | ETH Zurich, Zürich | ETH Zürich | Department Information Technology and Electrical Engineering | Research profile - ResearchGate, <https://www.researchgate.net/profile/Sebastian-Frey-8>
 53. A Parallel Ultra-Low Power Silent Speech Interface based on a Wearable, Fully-dry EMG Neckband The authors acknowledge support from the ETH-Domain Joint Initiative program (project UrbanTwin). - arXiv, <https://arxiv.org/html/2509.21964v1>
 54. INTERSPEECH.2024 - Others | Cool Papers - Immersive Paper Discovery, <https://papers.cool/venue/INTERSPEECH.2024?group=Others>
 55. High-resolution surface electromyographic activities of facial muscles during mimic movements in healthy adults: A prospective observational study - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC9790991/>
 56. Comparison of Facial Muscle Activation Patterns Between Healthy and Bell's Palsy Subjects Using High-Density Surface Electromyography - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC7835336/>
 57. An Improved Model for Voicing Silent Speech - ResearchGate, https://www.researchgate.net/publication/353491633_An_Improved_Model_for_Voicing_Silent_Speech
 58. Giusy Spacone - Google Scholar, <https://scholar.google.co.uk/citations?user=dGE8uMEAAAAJ&hl=th>